

When Classical Baselines Are Tuned as Carefully as the Quantum Model, Does Quantum Reservoir Computing Still Win?

Tushar Pandey

Texas A&M University

ORCID: 0000-0001-7448-5723

Abstract—Can a small quantum computer forecast a changing signal better than an ordinary classical method? Many studies say yes, but the classical baselines they compare against are often left untuned while the quantum model is carefully optimised. We ask what happens when the classical competitor is given the same care: the same size and the same tuning effort. Studying two popular reasons a quantum reservoir is thought to help, in exact simulations of small systems (up to eleven qubits) on two synthetic forecasting tasks, we find these two advantages disappear once the comparison is fair. Extra quantum measurements add nothing beyond a classical formula of the same size; and although a feedback loop turns a useless quantum setup into a working predictor, a well-tuned classical network still predicts slightly more accurately, by a statistically reliable margin. The point is not that quantum reservoirs can never win, but that two commonly cited advantages do not survive a fair comparison at this scale. We offer these matched comparisons as a reusable checklist for honest benchmarking; all results are reproducible from fixed random seeds.

Index Terms—quantum reservoir computing, fair baselines, time-series forecasting, echo-state networks, reproducibility

I. INTRODUCTION

Reservoir computing [1], [2] trains only a linear readout on top of a fixed nonlinear dynamical system, which makes it an attractive near-term application of quantum hardware: a fixed quantum reservoir avoids the trainability pathologies of variational circuits [3] while supplying a high-dimensional nonlinear feature map. Since the original proposals [4], [5], a growing body of work reports that quantum reservoir computing (QRC) outperforms classical reservoirs and recurrent networks on chaotic, financial, and biomedical time series [6], [7]. At the same time, dequantization and expressivity results [8], [9] caution that the function classes realised by such circuits can often be reproduced classically, which makes the choice of classical baseline decisive.

A recurring weakness in these reports is the classical baseline. When a quantum model is compared against an echo-state network (ESN) [1] or a polynomial readout, the strength of the conclusion depends entirely on whether the classical model was given an equal chance: the same feature-space capacity, the same hyperparameter-tuning budget, the same causal evaluation protocol, and a metric that does not implicitly favour the quantum side. In practice these conditions are often not met: baselines run at defaults while the quantum

model is tuned, and directional accuracy is reported without a persistence control.

This paper asks a narrow question: *when the classical baseline is matched in capacity and tuning, does the quantum reservoir’s reported advantage survive?* We study two representative advantage mechanisms in noiseless statevector simulation at small system size ($q \leq 11$) and find that on these tasks it does not. Our contribution is not the negative outcome per se but the two controls that produce it: a capacity-matched classical model and a tuning-matched classical recurrent network, together with a correctness-gated, deterministic simulation harness. We state the scope plainly: this is a simulation study at sizes that classical hardware can represent exactly ($q \leq 11$); the claim is “no advantage here, under fair comparison,” not “no advantage is possible.”

These two mechanisms admit the cleanest matched controls; other angles of the broader fair-comparison cartography are deferred to a companion treatment, to which the same methodology applies.

II. METHODS

A. Common harness and fairness controls

All experiments share a single evaluation harness with the following properties, several of which are the fairness controls whose absence we are critiquing.

a) *Causal, leakage-free evaluation*: Each series is split chronologically into training, validation, and test segments; every preprocessing step and hyperparameter is fit on the training/validation portion only and applied causally, and the test segment is touched once, for the final reported number.

b) *Validation-only tuning, applied equally*: All free hyperparameters (ridge penalty for every method, quantum evolution time τ , feedback gain k_{fb} , and the ESN spectral radius and leak rate) are selected by validation error. Importantly, the classical baselines receive the *same* tuning budget as the quantum model. This is the single most important control and the one most often missing.

c) *Matched feature dimension*: Where a quantum and a classical reservoir are compared, their readout feature dimensions are matched so that the comparison is not confounded by readout width.

d) *Correctness gates*: Before any result is generated, an automated self-test verifies state physicality ($\text{Tr} \rho = 1$, expectations in $[-1, 1]$), agreement of the two-point readout with the explicit operator expectation to machine precision, convergence of finite-shot estimates to the exact values, and a strictly causal feedback path that reproduces the open-loop reservoir at gain zero. Every reported number is read back from disk against a content hash; all results reproduce bit-for-bit.

e) *Metrics*: We report normalised root-mean-square error (NRMSE, normalised by the per-target standard deviation, so that $\text{NRMSE} = 1$ is the trivial mean predictor) and, for directional tasks, mean directional accuracy (MDA): the fraction of test points whose predicted next-move direction matches the observed one. MDA is quoted alongside the *persistence-MDA* of the naive predictor $\text{sign}(y_t - y_{t-h})$, the meaningful floor for a directional task, since an MDA above 0.5 is otherwise easy to obtain on trended data.

B. Quantum reservoirs

The quantum reservoir is a fully connected transverse-field Ising system,

$$H = \sum_{i < j} J_{ij} X_i X_j + v \sum_i Z_i, \quad J_{ij} \sim \mathcal{U}[0, 1], \quad v = 1, \quad (1)$$

fixed after random initialisation. Classical inputs are angle-encoded via single-qubit R_Y rotations; the state evolves for time τ ; and the readout consists of single-qubit expectations $\langle Z_i \rangle$ and, where indicated, two-point correlators $\langle Z_i Z_j \rangle$, which are estimable from the same computational-basis measurement record [10]. A recurrent variant carries memory across timesteps by partial-tracing the input qubits between encodings, following standard recurrent-QRC constructions [5]. Because Z_i and $Z_i Z_j$ are diagonal in the computational basis, all correlators are obtained from the state-diagonal in a single contraction, verified equal to the explicit operator expectation (Section II).

III. CASE STUDY I: A CORRELATOR READOUT ADDS NOTHING BEYOND A CAPACITY-MATCHED POLYNOMIAL

A. Setup

A frequently proposed route to quantum expressivity is to enrich the readout with two-point correlators $\langle Z_i Z_j \rangle$, growing the feature space quadratically; a fair comparison must ask whether these features carry information an explicit classical quadratic model does not already have.

We test this on a controlled testbed of $N = 5$ coupled Hénon maps: bounded, genuinely quadratic, multivariate dynamics with a tunable cross-asset coupling c . We compare three models on one-step prediction (all with the same windowed ridge readout and causal protocol): the tuned correlator QRC ($q = 8$, readout tuned over τ , encoding scale, and single- vs. two-reservoir ensemble); a degree-2 polynomial model (POLY2) whose features are all pairwise products of the windowed inputs, the *capacity-matched* classical control, given *more* features than the QRC; and the concatenation




Fig. 1. Case study I. One-step test NRMSE versus coupling c for the correlator QRC, POLY2, and their concatenation. The polynomial control sits below the QRC at every coupling and the concatenation tracks POLY2 almost exactly; the dotted line is the trivial-mean predictor ($\text{NRMSE} = 1$).

POLY2 $\oplus$ QRC, which tests directly whether the quantum features add anything to the polynomial basis. Observation noise of standard deviation 0.1 is added so that no model can solve the map exactly through its generating polynomial. The concatenation ablation is exactly the control omitted by typical hybrid-architecture reports.

B. Results

Table I reports the verified numbers, as mean $\pm$ standard deviation over five paired data-seed/reservoir-seed combinations per coupling. At the one-step horizon the POLY2 control beats the tuned QRC outright at every coupling, and at the strongest coupling the tuned QRC degrades past the trivial-mean predictor ($\text{NRMSE} > 1$) while the classical models remain predictive. Concatenating the quantum features onto POLY2 changes its error only within the seed-to-seed spread: the cat-vs-POLY2 relative reduction is only $\approx 2.6\%$ at every coupling, with a standard deviation across seeds comparable to its mean, the magnitude expected from a few extra mildly-regularising features rather than from complementary signal. The same verdict holds on mean directional accuracy: POLY2 reaches $\text{MDA} \approx 0.96$ ($\text{persistence} \leq 0.18$), the quantum model is no better, and the concatenation does not exceed POLY2. The coupling axis, which one might expect to favour the joint quantum encoding, instead makes the QRC relatively worse (Fig. 1).

TABLE I
CASE STUDY I (ONE-STEP, OBSERVATION NOISE 0.1): CORRELATOR QRC VS. THE CAPACITY-MATCHED POLY2 CONTROL AND THEIR CONCATENATION, ON COUPLED HÉNON DYNAMICS AT THREE COUPLINGS c . VALUES ARE MEAN $\pm$ S.D. OVER FIVE DATA-SEED $\times$ RESERVOIR-SEED PAIRS; LOWER NRMSE IS BETTER, $\text{NRMSE} < 1$ PREDICTIVE. “CAT” IS POLY2 $\oplus$ QRC; THE FINAL COLUMN IS THE RELATIVE NRMSE REDUCTION OF CAT OVER POLY2. WINDOWED FEATURE DIMENSIONS: QRC 72, POLY2 135, SO THE POLYNOMIAL IS GIVEN *more* CAPACITY, NOT LESS.

c	QRC	POLY2	cat	cat vs. POLY2
0.0	0.726 ± 0.026	0.363 ± 0.012	0.354 ± 0.012	$+2.6\% \pm 0.3$
0.1	0.766 ± 0.034	0.335 ± 0.012	0.326 ± 0.007	$+2.6\% \pm 1.5$
0.2	1.030 ± 0.258	0.305 ± 0.132	0.297 ± 0.125	$+2.7\% \pm 1.0$

We are precise about what this does and does not show. The coupled Hénon dynamics are exactly quadratic, so a degree-2 polynomial spans the true generating function class; this is by construction the most favourable possible setting for POLY2, and the result is correspondingly bounded. The claim we make is therefore not “a fair baseline always erases the edge,” but the narrower and fully supported one: *when the classical model already covers the data’s function class, a two-point correlator readout contributes nothing beyond it*. The correlator readout does help the QRC relative to a single-point readout, but that gain does not transfer into an advantage over a classical model of matched (indeed larger) capacity. A reader should read Case I as a clean refutation of the specific “correlators add quantum expressivity” argument on quadratic data, not as a universal statement. Two residual asymmetries bound the claim further: POLY2’s 135 features exceed the QRC’s 72, and the task’s function class is known in advance. The concatenation ablation is insensitive to the first (it asks whether the quantum features add anything to the polynomial basis, whatever that basis’s size), but an exactly equal-dimensional control and tasks whose governing function class is unknown a priori are the natural next step.

IV. CASE STUDY II: FEEDBACK IS A LIVE MECHANISM BUT LOSES TO A TUNING-MATCHED ESN

A. Setup

A distinct advantage claim concerns *feedback-driven QRC*, in which the measurement outcome at one step modulates the reservoir drive at the next, making the effective dynamics path-dependent rather than a fixed function of a finite input window. Path-dependence is something a fixed-window polynomial cannot reproduce, but it is precisely what a classical recurrent network also provides. The honest baseline for a feedback reservoir is therefore not a polynomial but a recurrent network, the echo-state network (ESN), and the test is whether *quantum* recurrence beats *classical* recurrence at equal tuning.

We use a nonstationary task built so that recurrence is necessary: a hidden two-state Markov process switches the sign of a first-order autoregressive coefficient, so each regime is individually predictable but the regime must be inferred online from the observable history. A correctness gate confirms the task genuinely requires regime inference (an oracle given the hidden state beats a fixed-window linear model, with the gap growing as the switching rate rises). We implement feedback as an additive, saturating term on the input drive angle, $\phi_t = x_t + k_{\text{fb}} \tanh(\bar{m}_{t-1})$, where $\bar{m}_{t-1}$ is the mean measurement at the previous step; $k_{\text{fb}} = 0$ recovers the open-loop reservoir exactly. We compare, over ten data seeds $\times$ ten reservoir seeds (100 paired runs; an initial 3×3 grid served only as a pilot), the tuned feedback QRC, its open-loop counterpart, a tuning-matched ESN, POLY2, and a linear model. For full disclosure of the tuning budgets: every method selects its ridge penalty over the six values $\{10^{-4}, 10^{-3}, \dots, 10\}$ on validation; the feedback QRC is tuned over $\tau \in \{1, 2\}$ and $k_{\text{fb}} \in \{0.5, 1, 2\}$ (six configurations, with the open-loop variant tuned over τ); the ESN is tuned over spectral radius $\{0.7, 0.9, 0.99\}$ and leak

rate $\{0.2, 0.5, 0.9\}$ (nine configurations, so the classical model receives *more* validation trials than the quantum one). The QRC uses $q = 5$ qubits (one input, four hidden) read out through its 15 $Z_i/Z_i Z_j$ observables over a five-step window (75 features); the ESN reservoir has 15 units so its windowed readout is also 75-dimensional. Both models therefore train exactly 75 ridge weights.

B. Results

Feedback is a genuine, live mechanism. In the 10×10 -seed comparison at $p_{\text{sw}} = 0.05$ (Table II), the open-loop reservoir is useless (NRMSE = 1.020 ± 0.036 , worse than the mean predictor) while closing the loop lifts it to 0.787 ± 0.059 , a 23% improvement on identical seeds; a single-seed sweep shows the same effect across $p_{\text{sw}} \in \{0.02, 0.05, 0.10\}$ (+10 to +19% gain; per-seed JSONs in the repository). This is the clearest case in our study of a quantum reservoir mechanism doing something its open-loop version cannot.

TABLE II
CASE STUDY II: THE FAIR TEST AT $p_{\text{sw}} = 0.05$ OVER 10×10 SEEDS (MEAN $\pm$ S.D.), ONE-STEP NRMSE AND MDA (PERSISTENCE MDA = 0.27). THE TUNING-MATCHED ESN BEATS THE FEEDBACK QRC; THE PAIRED NRMSE DIFFERENCE IS SIGNIFICANT (t -TEST $p \approx 10^{-21}$, WILCOXON $p \approx 10^{-16}$, BOOTSTRAP 95% CI [0.057, 0.078]) AND REMAINS SO AT THE DATA-SEED LEVEL ($p = 3.2 \times 10^{-3}$, HIERARCHICAL BOOTSTRAP CI [0.037, 0.100]).

method	NRMSE	MDA
ESN (tuned)	0.720 ± 0.060	0.737 ± 0.054
Linear	0.750 ± 0.056	0.735 ± 0.055
POLY2	0.762 ± 0.046	0.734 ± 0.055
Feedback QRC	0.787 ± 0.059	0.728 ± 0.051
Open-loop QRC	1.020 ± 0.036	0.740 ± 0.058

The mechanism being live, however, does not make it advantageous. Under the fair 10×10 -seed comparison (Table II), the tuning-matched ESN attains NRMSE = 0.720 ± 0.060 against the feedback QRC’s 0.787 ± 0.059 , an ESN advantage of 8.5%. Because the 100 runs are paired (shared data and reservoir seeds), we test the per-run difference directly: the mean paired difference is 0.067 NRMSE (ESN better), significant under a paired t -test ($t = 12.3$, $p \approx 1.4 \times 10^{-21}$), the Wilcoxon signed-rank test ($p \approx 9 \times 10^{-17}$), and a 20,000-resample bootstrap (95% CI [0.057, 0.078]), all taken across the 100 per-run differences. Runs sharing a data seed are correlated, however, so treating all 100 as independent risks pseudoreplication; we therefore also aggregate to the data-seed level. On the ten independent per-data-seed means the conclusion is unchanged: the ESN is ahead in 9 of 10 data seeds (paired $t = 3.99$, $p = 3.2 \times 10^{-3}$; Wilcoxon $p = 0.004$), and a hierarchical bootstrap that resamples data seeds and then reservoir seeds within them gives a 95% CI of [0.037, 0.100], still excluding zero. The ESN beats the feedback QRC in 92% of paired runs, and the QRC’s spread across realisations (s.d. = 0.059) matches the ESN’s (0.060), so the result is systematic rather than an unlucky draw. The feedback QRC in fact trails every classical baseline, including




Fig. 2. Case study II. One-step test NRMSE at $p_{sw} = 0.05$ (mean $\pm$ s.d., 10×10 seeds). Feedback rescues the open-loop reservoir (rightmost) to predictive performance, but the tuning-matched ESN and even a plain linear model stay ahead. Blue bars classical; dotted line trivial-mean.

a plain windowed linear model (0.750); note that on this task the quadratic POLY2 features (0.762) are themselves worse than the linear model, a reminder that added capacity is not automatically useful. On directional accuracy all models cluster near 0.73, far above the 0.27 persistence floor but statistically indistinguishable across methods (Fig. 2). Read as a bounded partial result rather than a simple loss: closing the loop carries a non-predictive reservoir to within 9% of the tuned ESN’s error, so quantum feedback is a real and nearly competitive recurrence mechanism; what it is not, on this task, is an advantage. Quantum feedback reconstructs the recurrent memory a classical recurrent network already provides, slightly less efficiently.

V. DISCUSSION

Two mechanisms, two representative advantage claims, one outcome: under a capacity-matched classical control in the first case and a tuning-matched classical recurrent control in the second, the quantum reservoir’s reported edge does not survive. We emphasise three points.

First, the controls are the contribution. The concatenation ablation ($POLY2 \oplus QRC$ vs. $POLY2$) directly measures whether quantum features add information to a classical model of matched capacity, and the same-budget ESN measures whether quantum recurrence beats classical recurrence on equal footing. Both are inexpensive and routinely omitted; including them changes the conclusion.

Second, “live” is not “advantageous.” The feedback reservoir does transform a useless open-loop model into a predictive one. A study that compared feedback QRC only against open-loop QRC, or against an untuned ESN, would have reported a success. The advantage evaporates only against a classical recurrent model tuned as hard as the quantum one.

Third, the scope is bounded and we state it as such. These are noiseless statevector simulations at $q \leq 11$, sizes a classical computer represents exactly, on synthetic tasks with low-dimensional classical inputs at a single one-step horizon; finite-shot and hardware-noise studies, longer horizons, and empirical series are natural extensions. We make no claim about larger systems, genuinely quantum-state inputs, or non-Clifford resources, where provable learning separations are

known to exist [11] and the relevant comparisons may behave differently. Our claim is the narrow, falsifiable one our experiments support: at currently simulable scales, with fair baselines, these two QRC advantage mechanisms do not beat their classical counterparts.

VI. CONCLUSION

We have shown that two representative quantum-reservoir advantage mechanisms, a two-point correlator readout and a feedback-driven reservoir, do not outperform capacity- and tuning-matched classical baselines on the tasks studied, in noiseless small-scale simulation. The controls behind this conclusion are inexpensive, decisive, and routinely omitted; we offer them as a reusable template for honest QRC benchmarking, not a verdict on quantum reservoir computing as a whole.

DATA AVAILABILITY

The result files and analysis scripts that reproduce every table and figure, including the per-seed JSONs for Case I and II, are at https://github.com/pandey-tushar/Quantum_Chaos_solver (branch `paper-7-cartography`); results reproduce bit-for-bit from the provided seeds.

ACKNOWLEDGMENT

The author used Claude (Anthropic) to assist with running and optimizing the simulation code. All technical content, results, and conclusions are the author’s own.

REFERENCES

- [1] H. Jaeger, “The ‘echo state’ approach to analysing and training recurrent neural networks,” GMD Technical Report 148, German National Research Center for Information Technology, 2001.
- [2] W. Maass, T. Natschlger, and H. Markram, “Real-time computing without stable states: A new framework for neural computation based on perturbations,” *Neural Comput.*, vol. 14, p. 2531, 2002.
- [3] J. R. McClean, S. Boixo, V. N. Smelyanskiy, R. Babbush, and H. Neven, “Barren plateaus in quantum neural network training landscapes,” *Nat. Commun.*, vol. 9, p. 4812, 2018.
- [4] K. Fujii and K. Nakajima, “Harnessing disordered-ensemble quantum dynamics for machine learning,” *Phys. Rev. Applied*, vol. 8, p. 024030, 2017.
- [5] K. Nakajima, K. Fujii, M. Negoro, K. Mitarai, and M. Kitagawa, “Boosting computational power through spatial multiplexing in quantum reservoir computing,” *Phys. Rev. Applied*, vol. 11, p. 034021, 2019.
- [6] P. Mujal, R. Martınez-Pena, J. Nokkala, J. Garcıa-Benı, G. L. Giorgi, M. C. Soriano, and R. Zambrini, “Opportunities in quantum reservoir computing and extreme learning machines,” *Adv. Quantum Technol.*, vol. 4, p. 2100027, 2021.
- [7] R. Martınez-Pena, G. L. Giorgi, J. Nokkala, M. C. Soriano, and R. Zambrini, “Dynamical phase transitions in quantum reservoir computing,” *Phys. Rev. Lett.*, vol. 127, p. 100502, 2021.
- [8] M. Schuld, R. Sweke, and J. J. Meyer, “Effect of data encoding on the expressive power of variational quantum-machine-learning models,” *Phys. Rev. A*, vol. 103, p. 032430, 2021.
- [9] R. Sweke, J.-P. Seifert, D. Hangleiter, and J. Eisert, “On the quantum versus classical learnability of discrete distributions,” *Quantum*, vol. 5, p. 417, 2021.
- [10] H.-Y. Huang, R. Kueng, and J. Preskill, “Predicting many properties of a quantum system from very few measurements,” *Nat. Phys.*, vol. 16, p. 1050, 2020.
- [11] H.-Y. Huang, M. Broughton, J. Cotler, S. Chen, J. Li, M. Mohseni, H. Neven, R. Babbush, R. Kueng, J. Preskill, and J. R. McClean, “Quantum advantage in learning from experiments,” *Science*, vol. 376, p. 1182, 2022.